Global AI Governance & Safety Guidelines 2025

[DEMONSTRATION SYNTHETIC DATASET - RESEARCH ETHICS POLICY]
Global AI Governance & Safety Guidelines 2025

Principle 1: Transparent Provenance and Attribution
AI-assisted retrieval systems must explicitly cite primary sources down to the document name and page number.
Hallucinated or fabricated citations constitute a critical violation of system integrity.

Principle 2: Groundedness and Abstention
If retrieved evidentiary context does not contain verifiable facts answering the user query,
the system must explicitly abstain by stating: 'The requested information was not found in the uploaded documents.'
Inventing plausible answers in high-stakes domains is strictly prohibited.

Principle 3: Privacy and Ephemeral Processing
User uploaded documents should never be persisted to external unsecured cloud services without explicit consent.
In-memory processing and local model inference (such as Ollama) are recommended for sensitive organizational data.